

Viet Nguyen

AI Data Engineer

Orange County, CA · (949) 992-7701 · vietdatapro@gmail.com · linkedin.com/in/viet-in-tech · github.com/viet-in-tech · viet-in-tech.github.io

SUMMARY

AI Data Engineer leading the enterprise-wide Amazon Quick rollout at ClinNEXUS, an AI-first healthcare startup, standardizing data sources and knowledge bases for AI consumption. Builds AI agents, RAG systems on Amazon Bedrock, end-to-end ML pipelines and containerized model APIs. Former robotics teacher who coached 4 consecutive VEX World Championship teams and explains technical ideas plainly to any audience. Pursuing an M.S. in Computer Science (AI) at Georgia Tech; AWS Certified Cloud Practitioner.

SKILLS

Generative AI & LLMs	Amazon Bedrock, Bedrock Knowledge Bases, RAG, LangChain, Groq (Llama 3.3), Prompt Engineering, Structured Outputs, AI Agents, LLM Evaluation, Titan Embeddings, OpenSearch Serverless, Hugging Face
Machine Learning	scikit-learn, TensorFlow/Keras, PyTorch, LightGBM, Random Forest, Gradient Boosting, Feature Engineering, Class Imbalance (SMOTE), GridSearchCV, ML Pipelines
Data & Analytics	SQL, Pandas, NumPy, PySpark, SciPy, Matplotlib, Seaborn, Plotly, Streamlit, EDA, Hypothesis & A/B Testing, Data Standardization & Schemas
AWS & Cloud	Amazon Quick (QuickSight), AWS Lambda, API Gateway, S3, EC2, EMR, CloudFront, IAM, VPC, CloudWatch, boto3, Vercel
MLOps & DevOps	Docker, Docker Hub, FastAPI, Pydantic, REST APIs, Terraform, Terraform Cloud, Git, GitHub
Languages	Python, SQL, TypeScript (Next.js), Go, Shell

EXPERIENCE

AI Data Engineer ClinNEXUS · Los Angeles, CA (Remote)	2026 — Present
- Own the enterprise-wide rollout and adoption of Amazon Quick, ensuring consistent deployment across all business units of an AI-first healthcare startup - Serve as subject matter expert on Amazon Quick capabilities, best practices, data methodologies and AI-driven workflows across the enterprise - Standardize and normalize datasets, data sources and knowledge bases for optimal AI consumption, partnering with cross-functional stakeholders to define data definitions, schemas and quality standards - Support population health, quality improvement and value-based care initiatives through AI-driven analytics, identifying automation opportunities across business functions	
Marketing Data Manager · Volunteer Association for Talent Development – Orange County (ATD-OC) · Anaheim, CA	2025 — Present
- Delivered a digital marketing report built in an AWS agentic AI workspace, with a 3-phase roadmap targeting 15% membership growth and 40% engagement improvement - Built and deployed an interactive marketing analytics dashboard with Claude and Vercel (<i>atdoc-dashboards.vercel.app</i>)	
Startup Co-Founder & Operations Lead TechX Robotics · Orange County, CA	2024 — 2026
- Managed full P&L for a \$100,000+ revenue robotics education business; audited 2 fiscal years, ran staff and contractor payroll, and calculated LLC partner distributions in Excel with Pivot Tables - Built an LLM-assisted compensation model, comparing parallel Claude, ChatGPT and Gemini outputs to evaluate owner's-distribution scenarios for 2 partners with differing work hours - Designed the competitive framework that led VEX Robotics teams to 1 World Championship	
Engineering & Robotics Teacher · VEX Robotics Coach Jeffrey Trail Middle School · Irvine, CA	2016 — 2026
- Taught robotics, AI, computer science and biotechnology; managed a \$15,000–\$20,000 annual engineering and robotics program budget - Coached teams to 4 consecutive World Championships and 5 consecutive State Championships	

I.T. & Ransomware Support Intern	2023
Manhattan Beachwear · Cypress, CA	
Supported a post-ransomware IT overhaul following a \$10M breach, deploying MFA, AI endpoint protection and secure remote access	
Loan Officer	2007 — 2008
Countrywide Home Loans · Palm Desert, CA	
Originated prime and subprime mortgage loans during the height of the U.S. mortgage crisis	

SELECTED PROJECTS

GENERATIVE AI & AGENTS	
ClassMaker.ai — AI Scheduling Agent · Next.js, TypeScript, Groq (Llama 3.3), Vercel	Apr 2026
Shipped at the Zero to Agent Global Hackathon: a full-stack agent that generates constraint-aware weekly class schedules as structured JSON within seconds and refines them through natural-language chat. Live at classmaker.ai.	
RAG System on Amazon Bedrock Knowledge Bases · Bedrock, Titan Embeddings, OpenSearch Serverless	Mar 2026
Built a RAG pipeline using the RetrieveAndGenerate API to ground LLM answers in private company documents.	
Stateful Chatbot with LangChain · LangChain, ChatBedrock	Mar 2026
Built a multi-turn chatbot with prompt templates, output parsers, conversation memory and CSV document injection.	
Serverless GenAI Flashcards App · AWS Lambda, API Gateway, Bedrock	Feb 2026
Connected Bedrock to a frontend via Lambda and API Gateway to turn study notes into flashcards, fully serverless.	
MACHINE LEARNING & MLOPS	
End-to-End ML on AWS · EC2, S3, IAM, VPC, scikit-learn	May 2026
Provisioned a VPC, EC2 and IAM role to train a diabetes classifier in the cloud, with data and model artifacts in S3.	
Medical Insurance Cost Prediction · TensorFlow/Keras, scikit-learn	Apr 2026
Built a deep learning regression pipeline (64→32 ReLU layers, Dropout, early stopping) that generalized cleanly with no overfitting.	
Containerized Survival Prediction API · Docker, FastAPI, scikit-learn	Apr 2026
Packaged a Random Forest model (81.56% test accuracy) and FastAPI service into a Docker image published to Docker Hub.	
House Price Prediction REST API · FastAPI, Pydantic, joblib	Mar 2026
Served a Random Forest regressor as a REST API with schema validation, startup model loading and a /health endpoint.	
Bank Customer Churn Prediction · scikit-learn, SMOTE, GridSearchCV	Jan 2026
Searched 108 hyperparameter combinations with feature engineering and SMOTE; F1 = 0.62 on a 20% minority class.	
DATA, CLOUD & RESEARCH	
AI Fairness in Mortgage Lending · Georgia Tech CS 6603	Nov 2024
Measured Disparate Impact and Statistical Parity on 2.5M Fannie Mae records; tested if Reweighting survives training.	
Big Data Benchmarking — Apache Spark on Amazon EMR · Georgia Tech CS 6675	Apr 2024
Benchmarked Spark and Hadoop across EC2 instance types; delivered cost/performance analysis from CloudWatch.	

EDUCATION

M.S. Computer Science (AI) · Georgia Tech	Exp. 2027
AI & Machine Learning Bootcamp · TripleTen	Exp. 2026
M.A. Teaching — Science Education · USC	2011
B.S. Chemistry · University of Texas at Arlington	

CERTIFICATION & TRAINING

AWS Certified Cloud Practitioner · AWS	2026
AI Engineering Program · SuperDataScience	2026
Reinforcement Learning for LLMs · RAGPACK.ai	2026
Supervised & Unsupervised ML · DeepLearning.AI	2025